

Diffusion Models

How AI Learns to Create Images from Noise

Presented by: *Anindya, Tamal and Din*

February 21, 2026

CSE 200: Technical Writing and Presentation

Bangladesh University of Engineering and Technology

Outline

Introduction

Core Idea

Forward Diffusion Process

Reverse Diffusion Process

Training

Insights

Applications

Limitations & Challenges

Conclusion

What Problem Are We Solving?



Real face



AI-generated face (*Sora*)

What Problem Are We Solving?



Real face



AI-generated face (*Sora*)

Generative Modeling

Can a machine learn to **create new, realistic samples** that are indistinguishable from real data?

What Problem Are We Solving?



Real face



AI-generated face (*Sora*)

Generative Modeling

Can a machine learn to **create new, realistic samples** that are indistinguishable from real data?

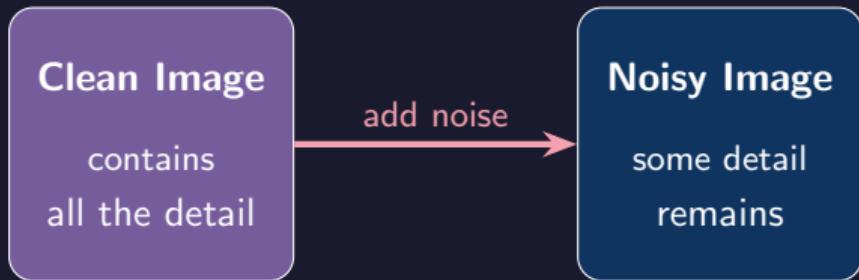
Diffusion models provide a stable, high-quality solution.

The Core Idea: Learning by Destroying

Clean Image

contains
all the detail

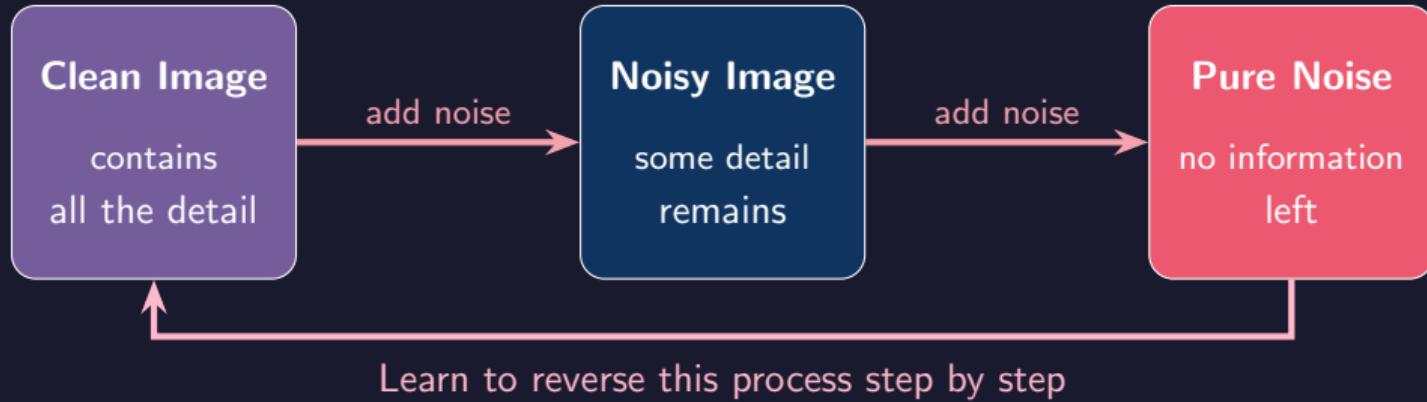
The Core Idea: Learning by Destroying



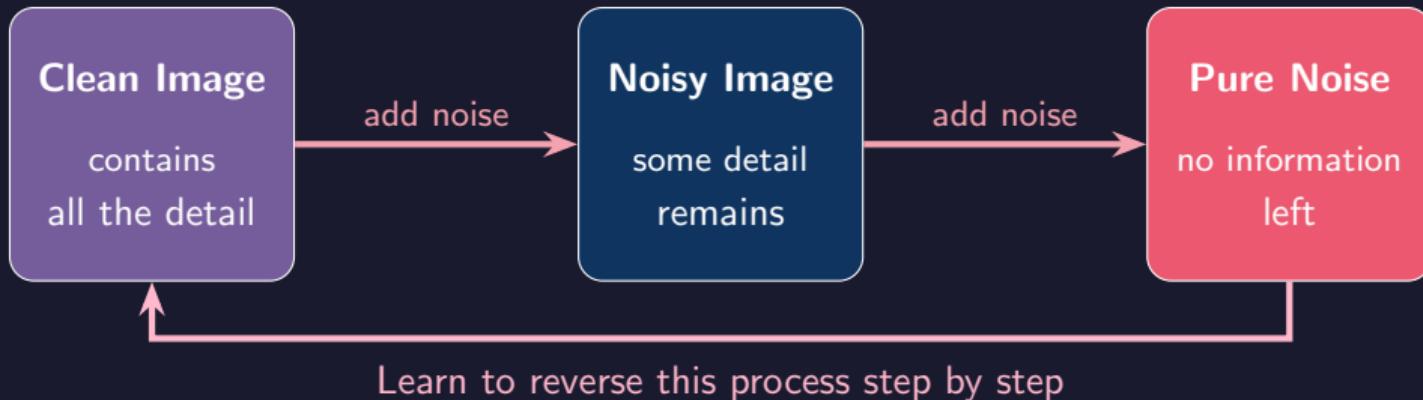
The Core Idea: Learning by Destroying



The Core Idea: Learning by Destroying



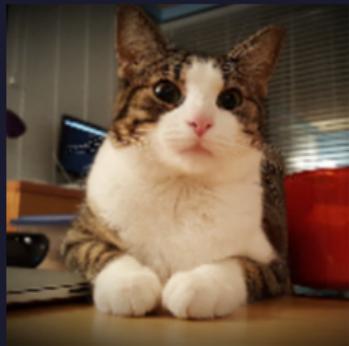
The Core Idea: Learning by Destroying



We Essentially Have Two Diffusion Processes

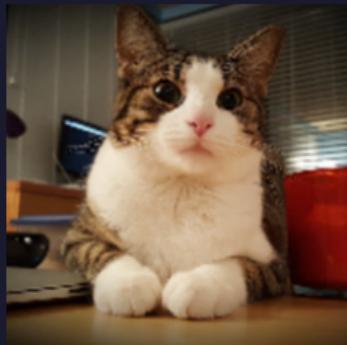
- The **Forward Diffusion Process** that gradually adds noise.
- The **Reverse Diffusion Process** that removes noise step by step.

The Forward Diffusion Process: Adding Noise



Clean Image

The Forward Diffusion Process: Adding Noise

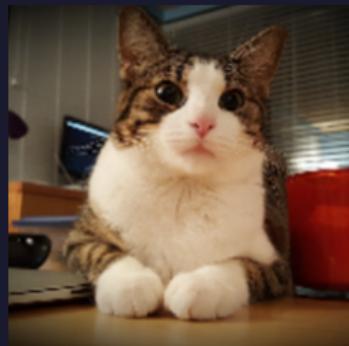


Clean Image



Slightly Noisy

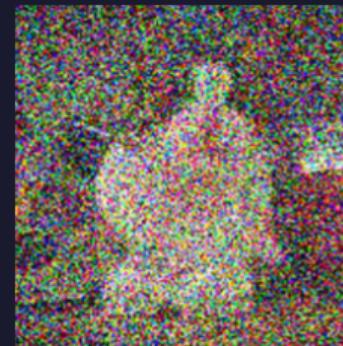
The Forward Diffusion Process: Adding Noise



Clean Image

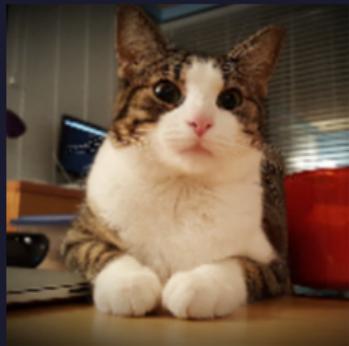


Slightly Noisy



Heavily Noisy

The Forward Diffusion Process: Adding Noise



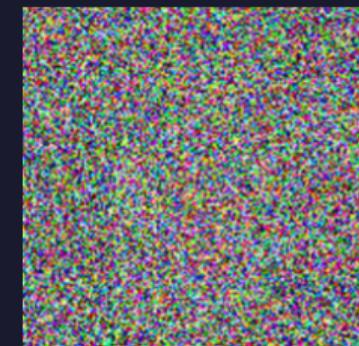
Clean Image



Slightly Noisy

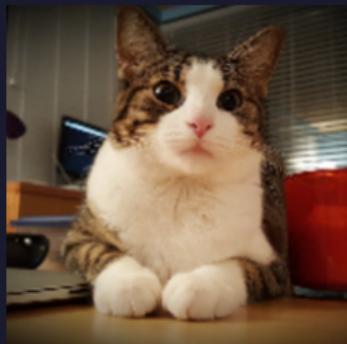


Heavily Noisy



Pure Noise

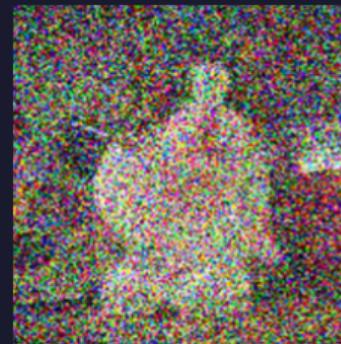
The Forward Diffusion Process: Adding Noise



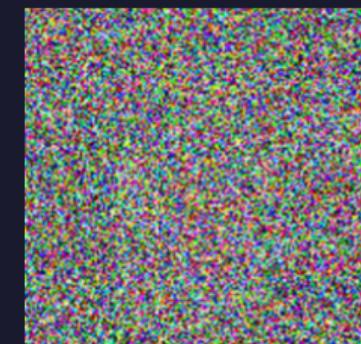
Clean Image



Slightly Noisy



Heavily Noisy

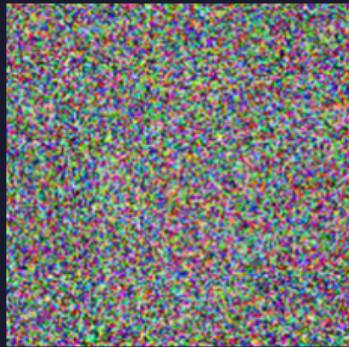


Pure Noise

What Happens

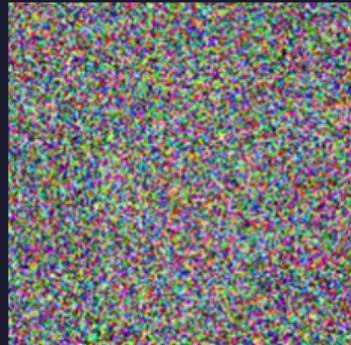
Take a real image and add a *tiny* bit of random noise. Repeat this hundreds of times until only noise remains.

The Reverse Diffusion Process: Removing Noise

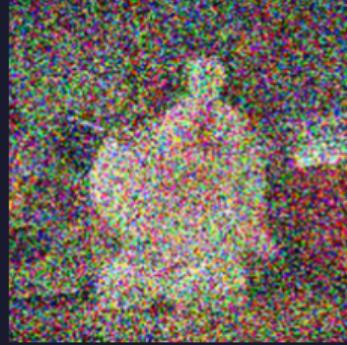


Pure Noise

The Reverse Diffusion Process: Removing Noise

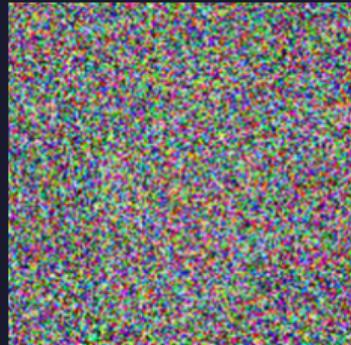


Pure Noise

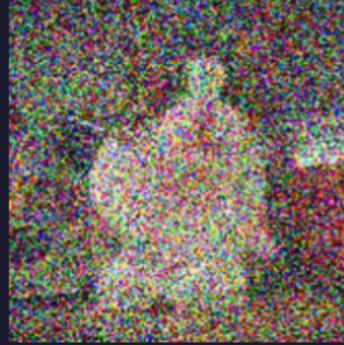


Slightly Less Noisy

The Reverse Diffusion Process: Removing Noise



Pure Noise



Slightly Less Noisy

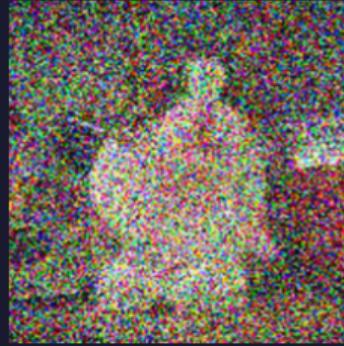


Almost Denoised

The Reverse Diffusion Process: Removing Noise



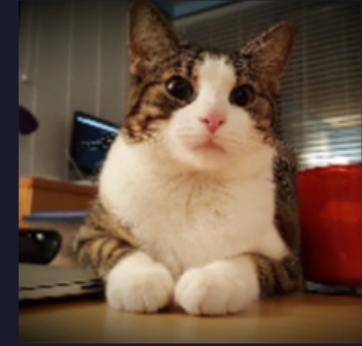
Pure Noise



Slightly Less Noisy

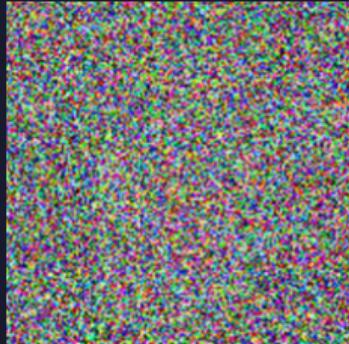


Almost Denoised

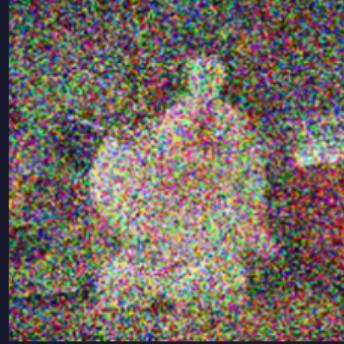


Fully Denoised

The Reverse Diffusion Process: Removing Noise



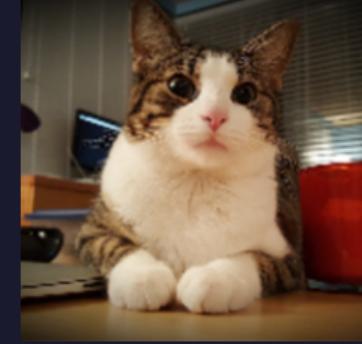
Pure Noise



Slightly Less Noisy



Almost Denoised



Fully Denoised

How Generation Works

- Start from **pure noise**
- Predict and remove a little noise at each step
- Repeat until a realistic image appears

Training

To produce realistic images, the model must learn to reverse the noise process.

Training

To produce realistic images, the model must learn to reverse the noise process.

But how does the model learn this reverse process?

How Do We Train It?

1. Take a real image from the training set

How Do We Train It?

1. Take a real image from the training set



2. Pick a random step and add the corresponding noise

How Do We Train It?

1. Take a real image from the training set



2. Pick a random step and add the corresponding noise



3. Ask the neural network: “*what noise was added?*”

How Do We Train It?

1. Take a real image from the training set



2. Pick a random step and add the corresponding noise



3. Ask the neural network: “*what noise was added?*”



4. Compare the prediction to the **actual noise**, minimize the error

Key Insight & Why It Works

Key Insight

- We add **known** noise
- So we have exact labels
- Training target is clear and reliable

Why It Works

- Model learns many small denoising steps
- Small steps are easier to learn
- Chaining them builds a full image

Why Diffusion Models Work So Well

Advantages

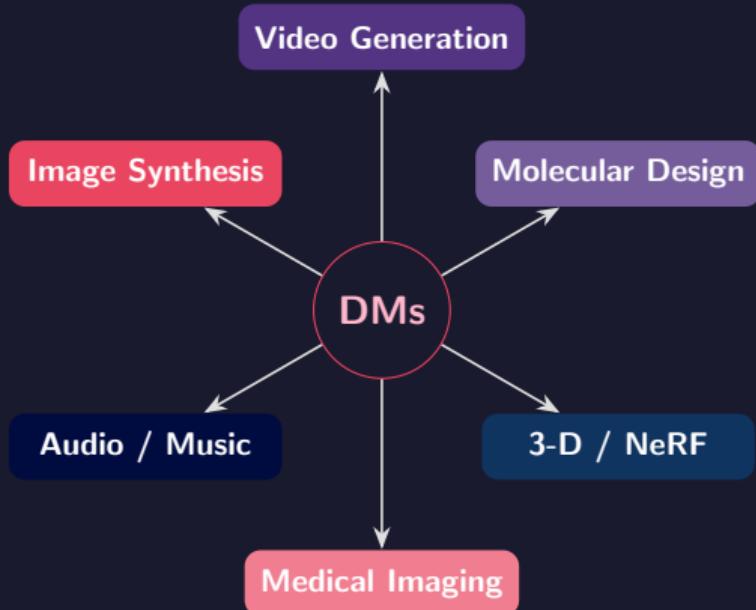
- **Stable training**
- **High Quality Samples**
- **Good Diversity** (less mode collapse than GANs)
- **Flexible Conditioning** (text-to-image, video-generation, audio synthesis)
- **Strong Theoretical Foundation**

In practice, diffusion models combine reliability with strong output quality.

Summary

Diffusion models generate data by learning to reverse a gradual noise process.

Applications



Some Real-World Tools Powered by Diffusion



Midjourney

Some Real-World Tools Powered by Diffusion



Midjourney



DALL·E

Some Real-World Tools Powered by Diffusion



Midjourney



DALL-E



OpenAI Sora

Limitations & Open Challenges

Current Limitations

- **Slow generation** — requires many denoising steps (hundreds to thousands)
- **High compute cost** — training requires significant GPU resources
- **Large model size** — not easy to run on consumer hardware

Open Questions

- Can we make generation **faster?** (Active research area)
- How do we ensure generated content is **safe and ethical?**
- Can we extend to **longer videos** and **interactive content?**

Thank You!

Thanks for Listening!

References

- [1] J. Ho et al. “**Denoising diffusion probabilistic models**”. In: *Proceedings of the 34th International Conference on Neural Information Processing Systems*. NIPS '20. Vancouver, BC, Canada: Curran Associates Inc., 2020.
- [2] Y. Song and S. Ermon. “**Generative modeling by estimating gradients of the data distribution**”. In: *Advances in neural information processing systems* 32 (2019).
- [3] J. Song et al. “**Denoising diffusion implicit models**”. In: *arXiv preprint arXiv:2010.02502* (2020).
- [4] A. Q. Nichol and P. Dhariwal. “**Improved denoising diffusion probabilistic models**”. In: *International conference on machine learning*. PMLR. 2021, pp. 8162–8171.