

# Diffusion Models

How AI Learns to Create Images from Noise

---

Presented by: *Anindya, Tamal and Din*

February 21, 2026

CSE 200: Technical Writing and Presentation

**Bangladesh University of Engineering and Technology**

# Outline

Introduction

Core Idea

Forward Diffusion Process

Reverse Diffusion Process

Training

Insights

Applications

Limitations & Challenges

Conclusion

# What Problem Are We Solving?



Real face




AI-generated face (*Sora*)

- Can a machine learn to **create new, realistic samples** that are *indistinguishable* from real data?
- In other words can it learn the underlying **probability distribution** of the data?

**Diffusion models** provide a *stable, high-quality* solution.


# The Core Idea: Learning by Destroying



## We Essentially Have Two Diffusion Processes

- The **Forward Diffusion Process** that gradually adds noise.
- The **Reverse Diffusion Process** that removes noise step by step.


# The Forward Diffusion Process: Adding Noise




Clean Image



Slightly Noisy



Heavily Noisy



Pure Noise


## What Happens

Take a real image and add a *tiny* bit of random noise. Repeat this hundreds of times until only noise remains.

# The Reverse Diffusion Process: Removing Noise




Pure Noise



Slightly Less Noisy



Almost Denoised



Fully Denoised

## How Generation Works

- Start from **pure noise**
- Predict and remove a little noise at each step
- Repeat until a realistic image appears

## Training

To produce realistic images, the model must **learn** to **reverse** the noise process.

**But how does the model learn this reverse process?**

## How Do We Train It?

1. Take a real image from the training set



2. Pick a random step and add the corresponding noise



3. Ask the neural network: “*what noise was added?*”



4. Compare the prediction to the **actual noise**, minimize the error

# Key Insight & Why It Works

## Key Insight

- We control the **noise process**
- Added noise is mathematically defined
- So the training signal is precise
- This makes the objective a well-defined *regression problem*

## Why It Works

- The task is simplified into **small, predictable steps**
- Each stem only removes a tiny amount of noise
- Repeating simple steps reconstructs complex structure

# Why Diffusion Models Work So Well

## Advantages

- **Stable training**
- **High Quality Samples**
- **Good Diversity** (less mode collapse than GANs)
- **Flexible Conditioning** (text-to-image, video-generation, audio synthesis)
- **Strong Theoretical Foundation** (probabilistic modeling, score matching)


In practice, diffusion models combine reliability with strong output quality.

## Summary

Diffusion models generate realistic data by learning the *underlying probability distribution* through reversing a *gradual, structured noise process*.

**But where is Diffusion Models actually applied?**

# Applications



# Some Real-World Tools Powered by Diffusion



Midjourney



DALL-E



OpenAI Sora

## Current Limitations

- **Slow generation:** requires many denoising steps (hundreds to thousands)
- **High compute cost:** training requires significant GPU resources
- **Large model size:** not easy to run on consumer hardware

## Open Challenges

- Can we make generation **faster**? (Active research area)
- How do we ensure generated content is **safe and ethical**?
- Can we extend to **longer videos** and **interactive content**?

# Thank You!

Thanks for Listening!

## References

- [1] J. Ho et al. “**Denoising diffusion probabilistic models**”. In: *Proceedings of the 34th International Conference on Neural Information Processing Systems*. NIPS '20. Vancouver, BC, Canada: Curran Associates Inc., 2020.
- [2] Y. Song and S. Ermon. “**Generative modeling by estimating gradients of the data distribution**”. In: *Advances in neural information processing systems* 32 (2019).
- [3] J. Song et al. “**Denoising diffusion implicit models**”. In: *arXiv preprint arXiv:2010.02502* (2020).
- [4] A. Q. Nichol and P. Dhariwal. “**Improved denoising diffusion probabilistic models**”. In: *International conference on machine learning*. PMLR. 2021, pp. 8162–8171.